

GLUU: BLIP2 Product Summary

Inference Architecture & Deployment Document

Project Overview

This document outlines the architecture, components, and deployment strategy for building a scalable, inference-based system using the BLIP2 model (Salesforce/blip2-flan-t5-xl) on AWS SageMaker, with a FastAPI backend hosted on ECS, API Gateway exposure, and OpenAI integration for final text product report.

This project will be executed by a designated team using the following resources:

- **Base Code:** `app.py` — core logic for BLIP2 execution and product report API.
- **Notebook Reference:** An official AWS sample notebook for BLIP2 inference using SageMaker.
- **Additional Resources:** Relevant AWS documentation, cost estimators, and architectural references.

Project Goal

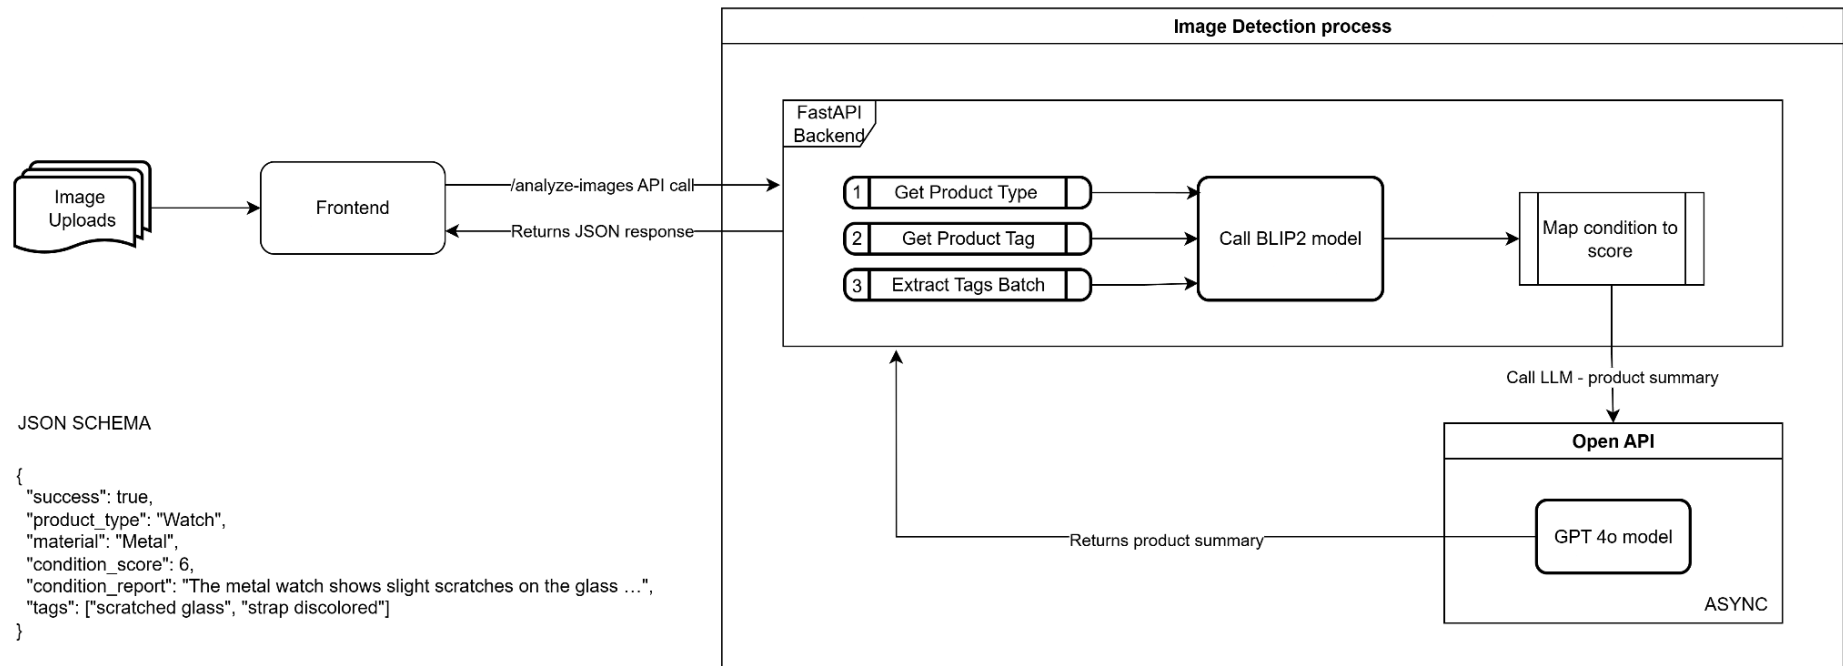
Implement the end-to-end architecture as described below, with a strong emphasis on:

- Reducing total predicted cost
- Finding cost-optimized alternatives without sacrificing core functionality

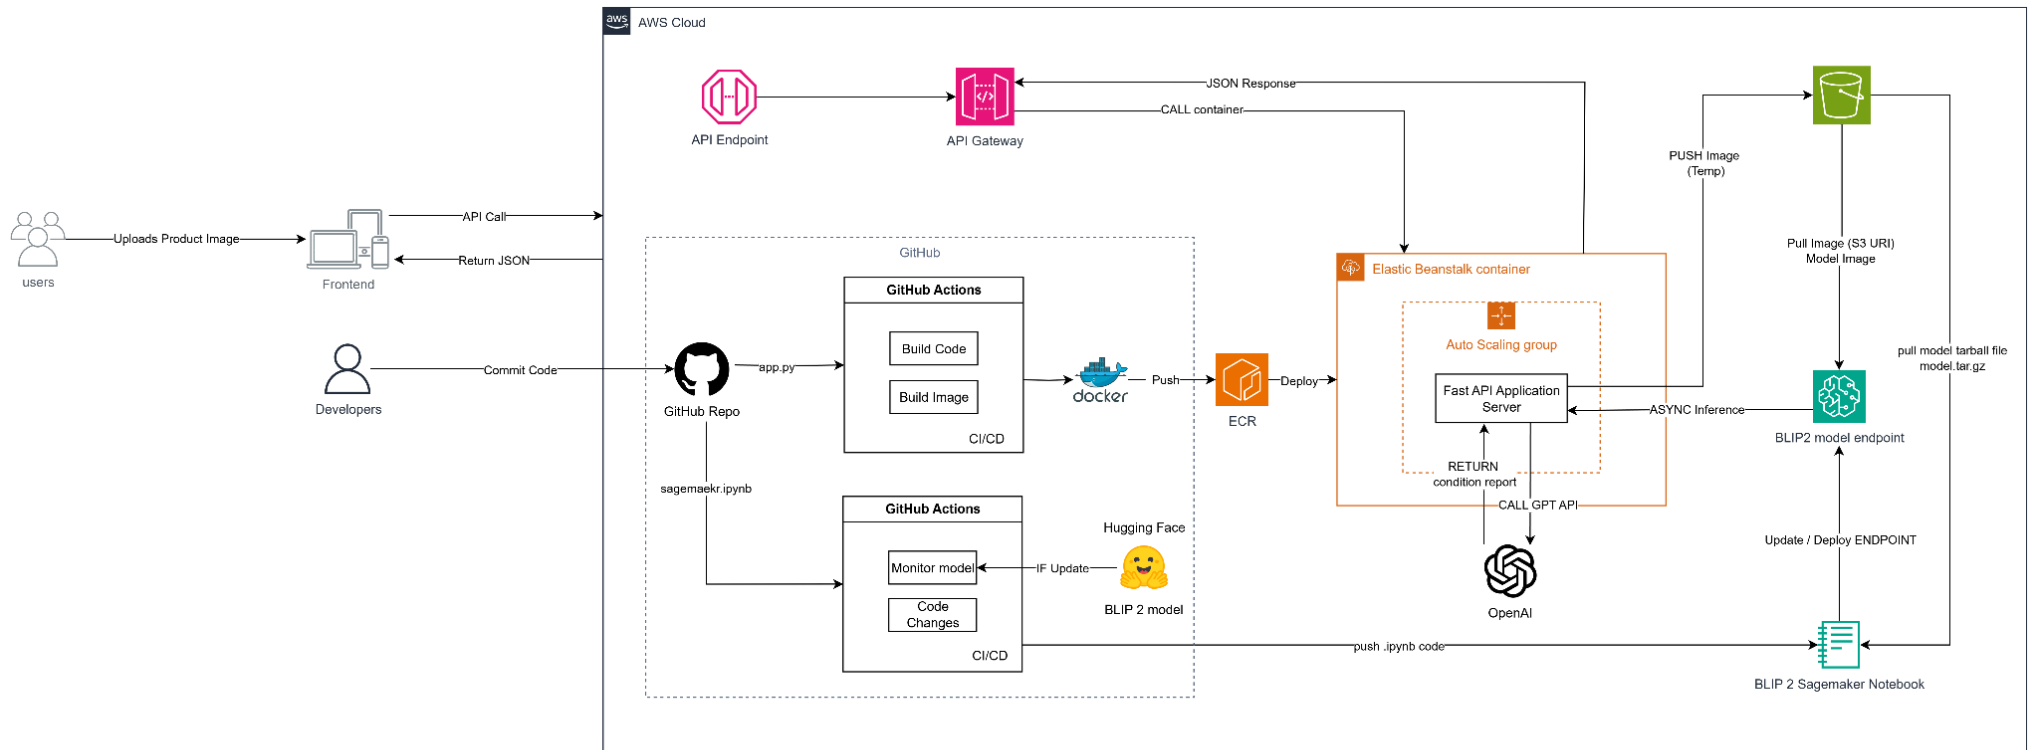
Leaving certain infrastructure choices **open-ended** for the technical team to decide during implementation, such as:

1. Whether to use a **Load Balancer or API Gateway** alone for routing
2. Whether to **host the backend on ECS (Fargate) or Elastic Beanstalk**
3. Finalizing the **async vs. sync inference** flow depending on performance testing
4. Whether **inference is executed through a SageMaker notebook** or via separate dedicated inference code in SageMaker Studio

System Overview



System Architecture (AWS)



Core Services (AWS and External)

Git & GitHub: Source code version control and collaboration

GitHub Actions: CI/CD automation for both backend and SageMaker deployment

Amazon ECS (Fargate) or Elastic Beanstalk: Hosts the API container in a serverless or managed environment (final choice to be determined during implementation)

Amazon ECR: Stores the container images

ALB + API Gateway: Handles HTTPS routing from client to ECS/Beanstalk (choice between ALB and direct API Gateway integration left open)

Route 53 + ACM (Optional): Custom domain and SSL management

Amazon S3: Temporary storage for uploaded images and model output (if async)

AWS SageMaker : Hosts the BLIP2 inference logic; may be executed directly via SageMaker notebook or via separate inference code in SageMaker Studio (final method open-ended)

OpenAI GPT-4: Generates product summaries from inference output

AWS Components

1. SageMaker Model Endpoint

Model: `Salesforce/blip2-flan-t5-xl`

Deployment Type: Real-time or Async Inference (Cost Saving)

Instance Type: `ml.g4dn.xlarge` (GPU)

Source Code:

- `inference.py`: Handles model loading and prediction
- `requirements.txt`: Contains dependencies (transformers, optimum, torch, etc.)
- Packaged as `blip2-endpoint.tar.gz` and uploaded to `s3://<bucket>/blip2/` via the sagemaker notebook code.

2. ECS FastAPI App (Backend API)

- Hosted in **ECS Fargate** with public **ALB (Application Load Balancer)** [Open Ended Requirement]
- Container built using `Dockerfile`
- Responsibilities:
 - Accept image uploads
 - Call SageMaker endpoint for inference
 - Format prompt and call OpenAI
 - Return structured JSON response

3. API Gateway

- Exposes a clean, secure REST/HTTP API to frontend or external apps
- Routes requests to ECS via ALB
- Can be enhanced with:
 - API Key
 - Lambda Authorizers or Cognito

4. OpenAI GPT-4o Integration

- Converts structured inference output into 50-word product summaries
- Model: `gpt-4o` or fallback to `gpt-4o-mini`, etc.

AWS Deployment Stages

STAGE 1: Setup SageMaker Studio & S3

- Launch Studio, create role with S3 access
- Use default or custom bucket
- Upload `blip2-endpoint.tar.gz` to `s3://<bucket>/blip2/`

STAGE 2: Deploy Model via Notebook or Inference Code File

Use `sagemaker.pytorch.PyTorchModel` and deploy to endpoint `blip2-endpoint`

STAGE 3: Build FastAPI App

- Code: `app.py`
- Handles file upload, calls SageMaker, formats OpenAI prompt
- Environment variables:
 - `SAGEMAKER_ENDPOINT=blip2-endpoint`
 - `OPENAI_API_KEY=<your-key>`

STAGE 4: Dockerize App

- `Dockerfile` and `requirements.txt`
- Build, tag, and push to Amazon ECR

STAGE 5: Deploy to ECS

- Create ECS cluster & Fargate service
- Attach ALB and open port 8080

STAGE 6: API Gateway Integration

- Create HTTP API with `ALB as integration (open ended)`
- Route: `POST /analyze`
- Optional: Connect Route 53 + ACM for HTTPS domain

Version Control and CI/CD

GitHub & GitHub Actions

All project code including: `app.py` (FastAPI backend), `Dockerfile`, `requirements.txt`, SageMaker notebook or inference model code is version controlled in GitHub via Git.

The CI/CD automation managed via **GitHub Actions** covers two primary components:

1. FastAPI Docker Backend (ECS Deployment)

On push to `main` or `release/*`:

1. Build Docker image from `app.py`
2. Push image to **Amazon ECR**
3. Trigger ECS service update to pull new image via AWS CLI/API

On changes to SageMaker notebook or inference code:

1. Automatic deployment to SageMaker using SageMaker Notebook Jobs or Papermill
2. Code review required before merge
3. If model updates are detected from Hugging Face, automatically retrain/redeploy the SageMaker endpoint

These pipelines ensure the entire backend and model logic remain version-controlled and automatically deployed.

SageMaker Endpoint Deployment

- The SageMaker deployment process is also automated through GitHub Actions.
- The SageMaker **notebook file** used to deploy the BLIP2 inference endpoint is version-controlled in GitHub.

CI/CD Steps:

Trigger:

- GitHub Actions is triggered on push/merge to `main`, `release/*`, or `sagemaker/*` branch
- Detects changes in the `sagemaker/` folder (including the deployment notebook)

Model Preparation in Notebook:

- The model (e.g. `Salesforce/blip2-flan-t5-xl`) is dynamically downloaded from Hugging Face within the notebook
- Inference logic and model code are bundled into a `.tar.gz` archive inside the notebook cell itself

Upload to S3:

- The notebook uploads the `blip2-endpoint.tar.gz` artifact to a configured S3 bucket (e.g., `s3://<bucket>/blip2/`)

Deploy/Update Endpoint (via notebook code):

- Using SageMaker SDK
- Notebook includes logic to **check for existing endpoint** and update it if necessary

Model Monitoring & Automation

- GitHub Actions periodically **checks for updates** to the model on Hugging Face using their API
- If a new version is found (e.g. SHA or version tag change):
 - Trigger the SageMaker notebook automatically using Amazon SageMaker Notebook Jobs or an nbconvert + CLI + Papermill-based execution
 - This allows auto-retraining or auto-redeployment based on upstream changes

Cost Considerations

Component	Estimated Monthly	Notes
SageMaker	~\$876	g4dn.xlarge - \$1.2/hr 24x7
ECS Fargate Task	~\$18	minimum 0.5vCPU, 1GB RAM full-time
OpenAI API (1k calls)	~\$6	\$0.006/call for gpt-4
API Gateway	Free or ~\$0.10	Under 1M reqs/month

Recommended - using async inference or cheaper OpenAI models for cost reduction.

Pricing Guidelines

SageMaker Pricing - <https://aws.amazon.com/sagemaker/pricing/>

ECS (Fargate) Pricing - <https://aws.amazon.com/fargate/pricing/>

API Gateway Pricing - <https://aws.amazon.com/api-gateway/pricing/>

Amazon S3 Pricing - <https://aws.amazon.com/s3/pricing/>

CloudWatch Logs Pricing - <https://aws.amazon.com/cloudwatch/pricing/>

Recommended Practices & Security

IAM Roles:

- ECS/Beanstalk task role to access S3 and invoke SageMaker
- SageMaker execution role for S3 input/output

Secrets: Use AWS Secrets Manager or SSM for storing OpenAI API keys

Networking: Restrict access with security groups, VPCs, and HTTPS-only APIs

Lifecycle Policies: Auto-delete S3 images after N days if async is used

Project Resources

Gluu GitHub Repo: <https://github.com/gluudevteam/gluu-image-detection-blip2>

AWS SageMaker Documentation:

<https://docs.aws.amazon.com/pdfs/sagemaker/latest/dg/sagemaker-dg.pdf>

AWS Docker Deployment:

<https://aws.amazon.com/blogs/devops/deploy-a-docker-application-on-aws-elastic-beanstalk-with-gitlab/>

AWS API Documentation:

- <https://docs.aws.amazon.com/apigateway/latest/developerguide/how-to-deploy-api.html>
- <https://docs.aws.amazon.com/apigateway/latest/developerguide/welcome.html>

AWS Cost Optimization:

- Async Inference:
<https://docs.aws.amazon.com/sagemaker/latest/dg/async-inference.html>
- Multi-model Endpoints:
<https://docs.aws.amazon.com/sagemaker/latest/dg/multi-model-endpoints.html>
- SageMaker Auto Shutdown:
<https://github.com/aws-samples/sagemaker-studio-auto-shutdown-extension>

AWS Sample Notebooks:

- BLIP2 on SageMaker:
<https://github.com/aws-samples/amazon-sagemaker-genai-content-moderation/blob/main/blip2-sagemaker.ipynb>
- Gen AI on SageMaker:
<https://github.com/aws-samples/sagemaker-genai-hosting-examples/blob/main/Llama2/Llama2-7b/LMI/llama2-7b.ipynb>

Open AI Documentation:

- <https://openai.github.io/openai-agents-python/>
- <https://platform.openai.com/docs/models>

YouTube Playlist for SageMaker:

<https://youtube.com/playlist?list=PLnJDJmQkmJTc2NgcqRzDuhcQqXZTSxxOm>

Extra Reading:

<https://medium.com/@himanshusangshetty/deploying-a-ml-model-on-amazon-sagemaker-421f3510fb8f>